

Lab 4: Creating Simple MDP MATLAB Environment with a Q-Learning Agent

Abhishek M J - CS21B2018

21-02-2024

MDP Matlab Environment

```
MDP = createMDP(8, ["left"; "right"])
```

- This line creates a Markov Decision Process (MDP) object using the createMDP function.
- The first argument, 8, specifies the number of states in the MDP. Imagine eight distinct positions or situations in your environment.
- The second argument, ["left"; "right"], defines the possible actions that can be taken in each state. In this case, you can either move "left" or "right".

MDP =

GenericMDP with properties:

```
CurrentState: "s1"  
States: [8×1 string]  
Actions: [2×1 string]  
T: [8×8×2 double]  
R: [8×8×2 double]  
TerminalStates: [0×1 string]  
ProbabilityTolerance: 8.8818e-16
```

MDP.States

- It displays the numbers from 1 to 8, representing the unique identifiers for each state.

ans =

8×1 string array

```
"s1"  
"s2"  
"s3"  
"s4"  
"s5"  
"s6"  
"s7"  
"s8"
```

MDP.Actions

- It shows “left” and “right”, the two actions you can take from any state.

ans =

2×1 string array

```
"left"  
"right"
```

MDP.T

- Each element in the array MDP.T(s, a, s') tells you the probability of transitioning from state s to state s' when action a is taken.

ans(:, :, 1) =

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

ans(:, :, 2) =

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

MDP.R

- Each element in the array MDP.R(s, a) tells you the reward you receive immediately after taking action a in state s.

```
ans(:,:,1) =
```

```

10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10

```

```
ans(:,:,2) =
```

```

10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10
10    -1    -1    -1    -1    -1    -1    10

```

Defining Rewards and Transitions Probabilities

```
MDP.TerminalStates = ["s1";"s8"];
```

- This line specifies that states 1 and 8 are terminal states. Once the agent reaches one of these states, the episode ends.

```
nS = numel(MDP.States);
```

```
nA = numel(MDP.Actions);
```

- This line calculates the number of states and actions in the MDP.

```
MDP.R = -1*ones(nS,nA);
```

- This line initializes the reward array with -1 for all state-action pairs.

```
MDP.R(:, state2idx(MDP, MDP.TerminalStates), :) = 10;
```

- This line sets the reward to 10 for all state-action pairs that lead to a terminal state.

```
MDP.T(1, 1, 1) = 1;
```

```
MDP.T(1, 2, 2) = 1;
```

```
MDP.T(2, 1, 1) = 1;
```

```
MDP.T(2, 3, 2) = 1;
```

```
MDP.T(3, 2, 1) = 1;
```

```
MDP.T(3, 4, 2) = 1;
```

```
MDP.T(4, 3, 1) = 1;
```

```
MDP.T(4, 5, 2) = 1;
```

```
MDP.T(5, 4, 1) = 1;
```

```
MDP.T(5, 6, 2) = 1;
```

```
MDP.T(6, 5, 1) = 1;
```

```
MDP.T(6, 7, 2) = 1;
```

```
MDP.T(7, 6, 1) = 1;
```

```
MDP.T(7, 8, 2) = 1;
```

```
MDP.T(8, 7, 1) = 1;
```

```
MDP.T(8, 8, 2) = 1;
```

- This block of code sets the transition probabilities for each state-action pair. For example, `MDP.T(1, 2, 2) = 1` sets the probability of transitioning from state 1 to state 2 when action 2 (“right”) is taken to 1.
- In simpler terms, the agent moves to right from state 1 it will always (with probability 1) reach state 2.

```
MDP.T
```

```
MDP.R
```

```
env = rlMDPEnv(MDP)
```

- This line creates a reinforcement learning environment using the `rlMDPEnv` function. The environment is based on the MDP object we created earlier.

Define Q-Table and Initialize Agent

```
state_information = getObservationInfo(env)
```

```
action_information = getActionInfo(env)
```

- This line retrieves the observation and action information from the environment.

```
state_information =
```

[rlFiniteSetSpec](#) with properties:

```
    Elements: [8×1 double]
        Name: "MDP Observations"
Description: [0×0 string]
    Dimension: [1 1]
    DataType: "double"
```

```
action_information =
```

[rlFiniteSetSpec](#) with properties:

```
    Elements: [2×1 double]
        Name: "MDP Actions"
Description: [0×0 string]
    Dimension: [1 1]
    DataType: "double"
```

```
qTable = rlTable(state_information, action_information)
qTable.Table
```

- This line creates a Q-table using the `rlTable` function. The Q-table is a matrix that stores the Q-values for each state-action pair.

```
qTable =
```

[rlTable](#) with properties:

```
    Table: [8×2 double]
```

```
ans =
```

```
    0    0
    0    0
    0    0
    0    0
    0    0
    0    0
    0    0
    0    0
```

```
qTable.Table = ones(size(qTable.Table))*5
qTable.Table
```

- This line initializes the Q-table with a constant value of 5.

`qTable =`

`rlTable` with properties:

Table: [8×2 double]

`ans =`

5	5
5	5
5	5
5	5
5	5
5	5
5	5
5	5

```
qRepresentation = rlQValueRepresentation(qTable, state_information, action_information)
qRepresentation.Options
qRepresentation.Options.L2RegularizationFactor = 0;
qRepresentation.Options.LearnRate = 0.01;
```

- This block of code creates a Q-value representation using the `rlQValueRepresentation` function. The Q-value representation is used to define the learning parameters for the Q-learning agent.
- The `L2RegularizationFactor` and `LearnRate` options are set to 0 and 0.01, respectively.

```
qRepresentation =
```

[rlQValueRepresentation](#) with properties:

```
    ActionInfo: [1×1 rl.util.rlFiniteSetSpec]  
    ObservationInfo: [1×1 rl.util.rlFiniteSetSpec]  
    Options: [1×1 rl.option.rlRepresentationOptions]
```

```
ans =
```

[rlRepresentationOptions](#) with properties:

```
    LearnRate: 0.0100  
    GradientThreshold: Inf  
    GradientThresholdMethod: "l2norm"  
    L2RegularizationFactor: 1.0000e-04  
    UseDevice: "cpu"  
    Optimizer: "adam"  
    OptimizerParameters: [1×1 rl.option.OptimizerParameters]
```

```
agentOpts = rlQAgentOptions  
agentOpts.EpsilonGreedyExploration  
agentOpts.EpsilonGreedyExploration.EpsilonDecay = 0.01  
qAgent = rlQAgent(qRepresentation, agentOpts)
```

- This block of code creates a Q-learning agent using the `rlQAgent` function. The agent uses the Q-value representation and agent options we defined earlier.
- The `EpsilonGreedyExploration.EpsilonDecay` option is set to 0.01, which means the exploration rate decreases by 0.01 after each episode.
- The exploration rate determines the probability of the agent taking a random action instead of the action with the highest Q-value.

```
agentOpts =
```

```
rlQAgentOptions with properties:
```

```
    SampleTime: 1
    DiscountFactor: 0.9900
    EpsilonGreedyExploration: [1×1 rl.option.EpsilonGreedyExploration]
    CriticOptimizerOptions: [1×1 rl.option.rlOptimizerOptions]
    InfoToSave: [1×1 struct]
```

```
ans =
```

```
EpsilonGreedyExploration with properties:
```

```
    EpsilonDecay: 0.0050
    Epsilon: 1
    EpsilonMin: 0.0100
```

```
agentOpts =
```

```
rlQAgentOptions with properties:
```

```
    SampleTime: 1
    DiscountFactor: 0.9900
    EpsilonGreedyExploration: [1×1 rl.option.EpsilonGreedyExploration]
    CriticOptimizerOptions: [1×1 rl.option.rlOptimizerOptions]
    InfoToSave: [1×1 struct]
```

```
qAgent =
```

```
rlQAgent with properties:
```

```
    AgentOptions: [1×1 rl.option.rlQAgentOptions]
    UseExplorationPolicy: 0
    ObservationInfo: [1×1 rl.util.rlFiniteSetSpec]
    ActionInfo: [1×1 rl.util.rlFiniteSetSpec]
    SampleTime: 1
```

Train the Q-Learning Agent

```
trainOpts = rlTrainingOptions
trainOpts.MaxStepsPerEpisode = 10;
trainOpts.MaxEpisodes = 100;
trainOpts.StopTrainingCriteria = "AverageReward";
```



```
trainOpts.StopTrainingValue = 13;
trainOpts.ScoreAveragingWindowLength = 30;
```

- This block of code creates training options using the `rlTrainingOptions` function. The training options specify the maximum number of steps and episodes, as well as the stopping criteria for training.

```
trainOpts =
```

`rlTrainingOptions` with properties:

```

        MaxEpisodes: 500
        MaxStepsPerEpisode: 500
        StopOnError: "on"
ScoreAveragingWindowLength: 5
        StopTrainingCriteria: "AverageSteps"
        StopTrainingValue: 500
        SaveAgentCriteria: "none"
        SaveAgentValue: "none"
        SaveAgentDirectory: "savedAgents"
        Verbose: 0
        Plots: "training-progress"
        UseParallel: 0
ParallelizationOptions: [1x1 rl.option.ParallelTraining]
```

```

5      5
5      5
5      5
5      5
5      5
5      5
5      5
5      5
```

```
QTable0 = getLearnableParameters(getCritic(qAgent));
disp(QTable0{1})
```

- This line retrieves the initial Q-table from the Q-learning agent.

Data =

struct with fields:

```
    Observation: [1×1 struct]
      Action: [1×1 struct]
      Reward: [1×1 timeseries]
      IsDone: [1×1 timeseries]
SimulationInfo: [1×1 struct]
```

cumulativeReward =

9

```
doTraining = true;
if doTraining
    trainingStats = train(qAgent, env, trainOpts);
else
    load('genericMDPQAgent.mat', 'qAgent')
end
```

- This block of code trains the Q-learning agent using the train function. The training process is controlled by the training options we defined earlier.
- The trainingStats variable stores the training statistics, such as the average reward per episode and the total number of steps taken.

ans =

8×2 single matrix

5.0000	5.0000
6.1872	4.6774
4.7846	4.5259
4.3065	4.2784
4.2419	4.2593
4.5215	4.8004
4.6732	6.2947
5.0000	5.0000

```
Data = sim(qAgent, env)
cumulativeReward = sum(Data.Reward)
```

- This line simulates the Q-learning agent in the environment using the sim function. The Data variable stores the observations, actions, and rewards collected during the simulation.
- The cumulativeReward variable calculates the total reward obtained by the agent during the simulation.

```
QTable1 = getLearnableParameters(getCritic(qAgent));
disp(QTable1{1})
```

- This line retrieves the final Q-table from the Q-learning agent after training.

